



L2D


Reinforcement learning

Dr Neythen Treloar

Resources




Gym is a standard API for reinforcement learning, and a diverse collection of reference environments



The Gym interface is simple, pythonic, and capable of representing general RL problems:

```
import gym
env = gym.make("LunarLander-v2", render_mode="human")
observation, info = env.reset(seed=42)
for _ in range(1000):
    action = policy(observation) # User-defined policy function
    observation, reward, terminated, truncated, info = env.step(action)

    if terminated or truncated:
        observation, info = env.reset()
env.close()
```



[Docs](#) » Welcome to Spinning Up in Deep RL!

[Edit on GitHub](#)

Welcome to Spinning Up in Deep RL!



User Documentation

Available free here:

<https://web.stanford.edu/class/psych209/Readings/SuttonBartoIPRLBook2ndEd.pdf>

<https://github.com/Farama-Foundation/Gymnasium>

<https://spinningup.openai.com/en/latest/>

Deep reinforcement learning for the control of microbial co-cultures in bioreactors

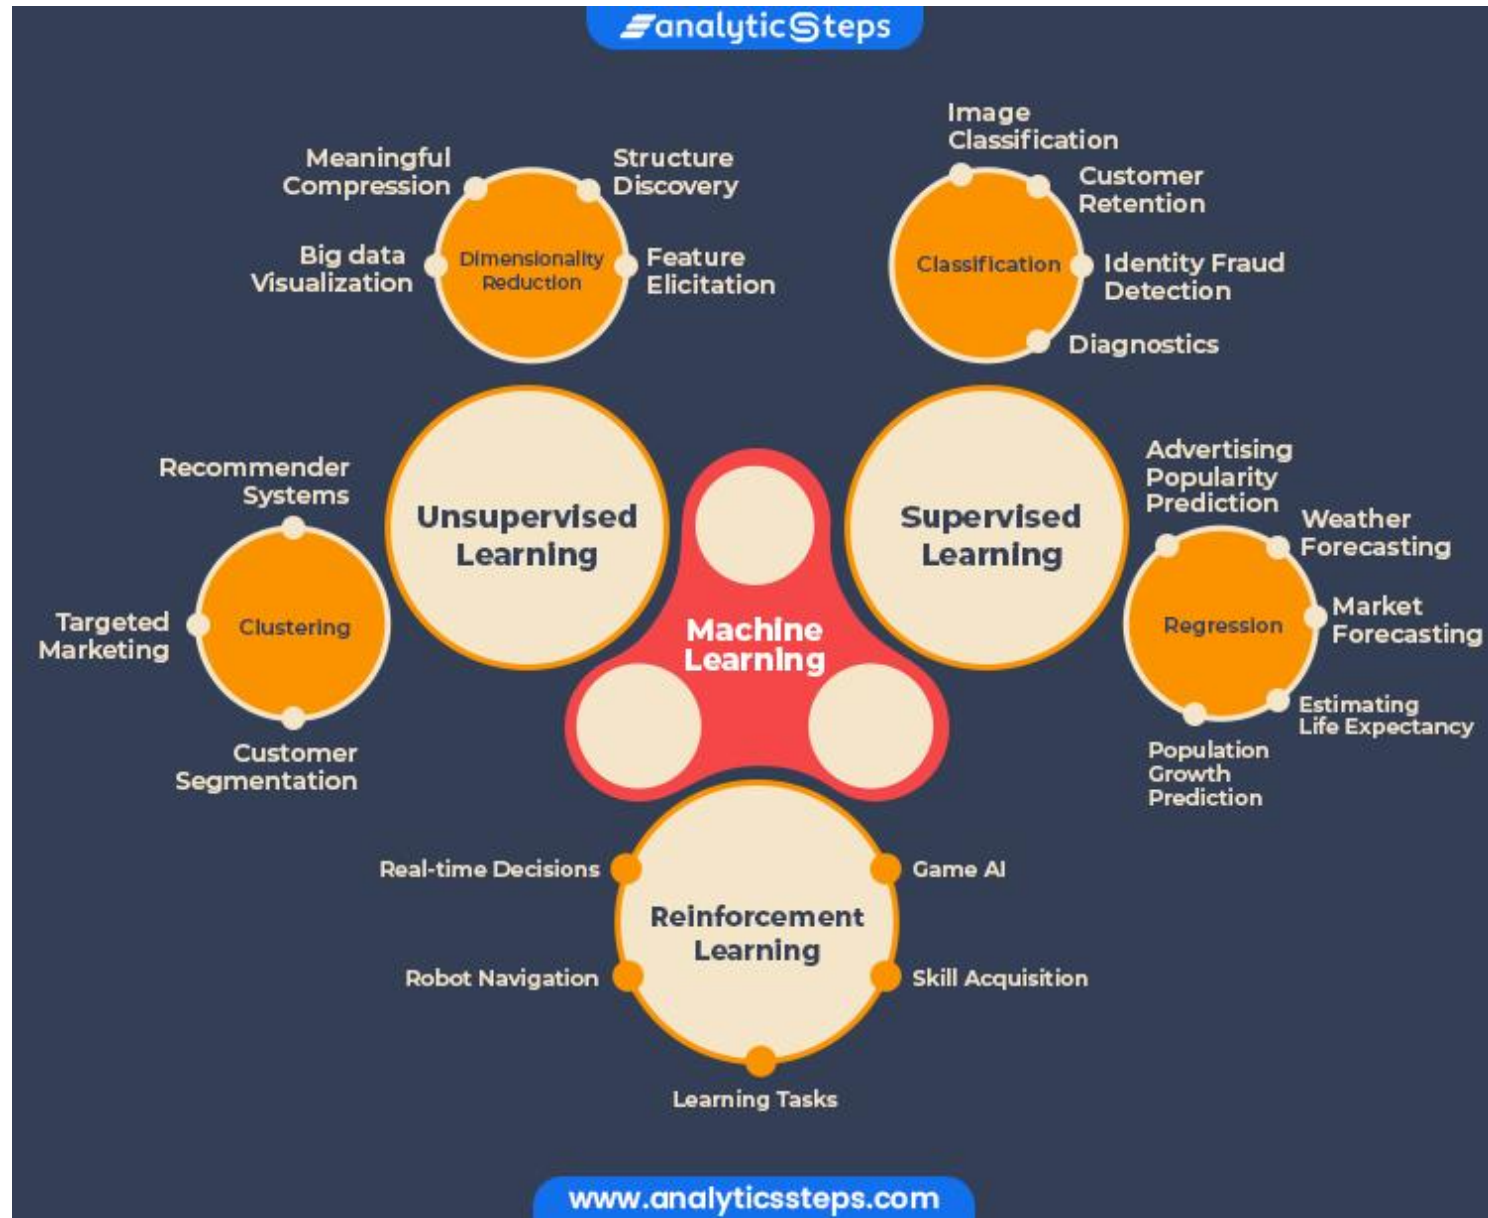
Neythen J. Treloar, Alex J. H. Fedorec, Brian Ingalls , Chris P. Barnes 

<https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1007783>

Brief overview

- Session 1: RL intro and tabular methods
 - Monte Carlo techniques
 - Temporal difference
 - Q-learning (off policy)
 - SARSA (on policy)
- Session 2: Deep reinforcement learning
 - Difficulties using neural networks for RL and the solutions

Introduction to machine learning approaches



Reinforcement learning demos

Learning good behaviour through trial and error

https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_dp.html

The ancient history

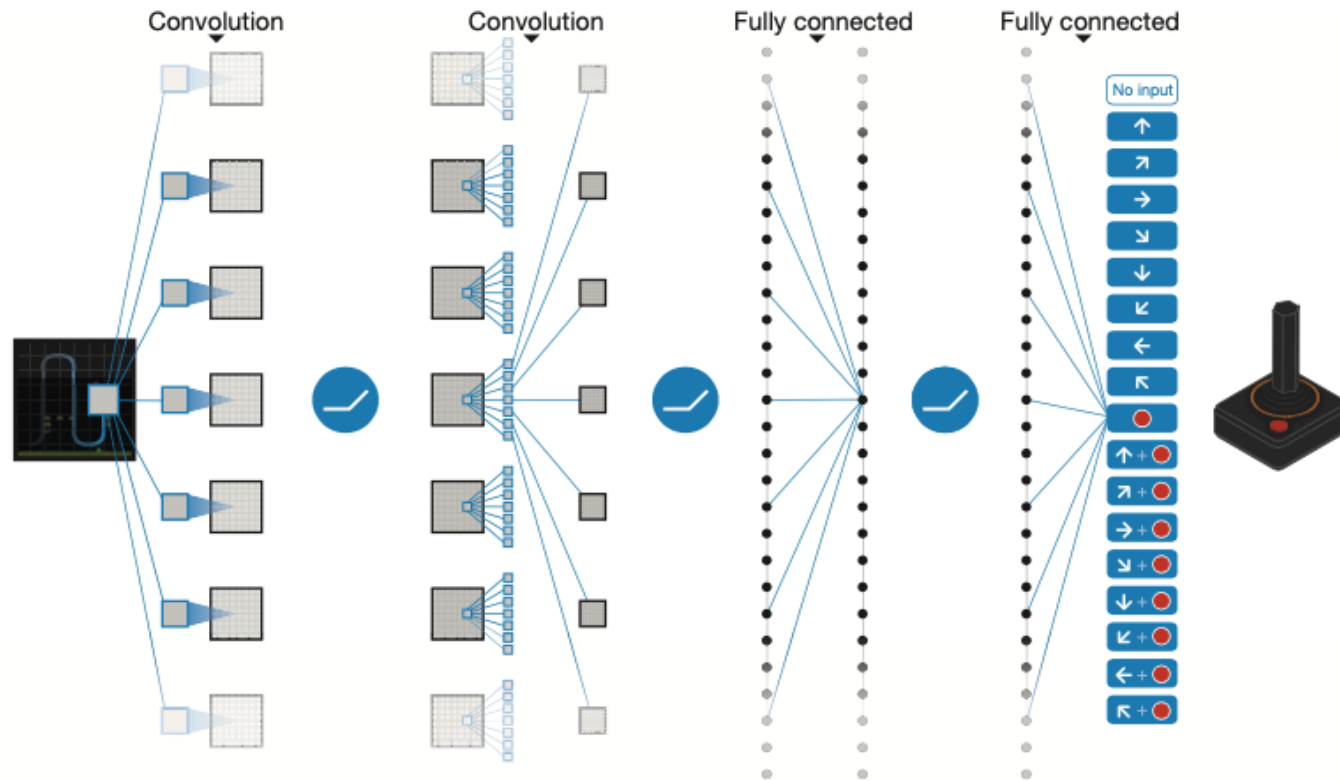
- Has roots in both psychology and control theory
- Control:
 - Bellman developed dynamic programming to solve the Bellman equations
 - The only feasible way to solve general stochastic optimal control problems
- Trial and error learning in animal behaviour:
 - How do animals learn behaviour that maximises positive effects
 - Pavlov etc.
- Development of Q-learning by Chris Watkins 1989 [1]
- TD-gammon, backgammon playing program 1994 [2]
- Page 13 in Sutton and Barto

[1] Watkins, C.J.C.H., 1989. Learning from delayed rewards

[2] Tesauro, G., 1994. TD-Gammon, a self-teaching backgammon program, achieves master-level play. *Neural computation*

Modern history

- Development of DQN, Q learning with neural networks (2015)



Mnih, V et al. 2015. Human-level control through deep reinforcement learning. *nature*

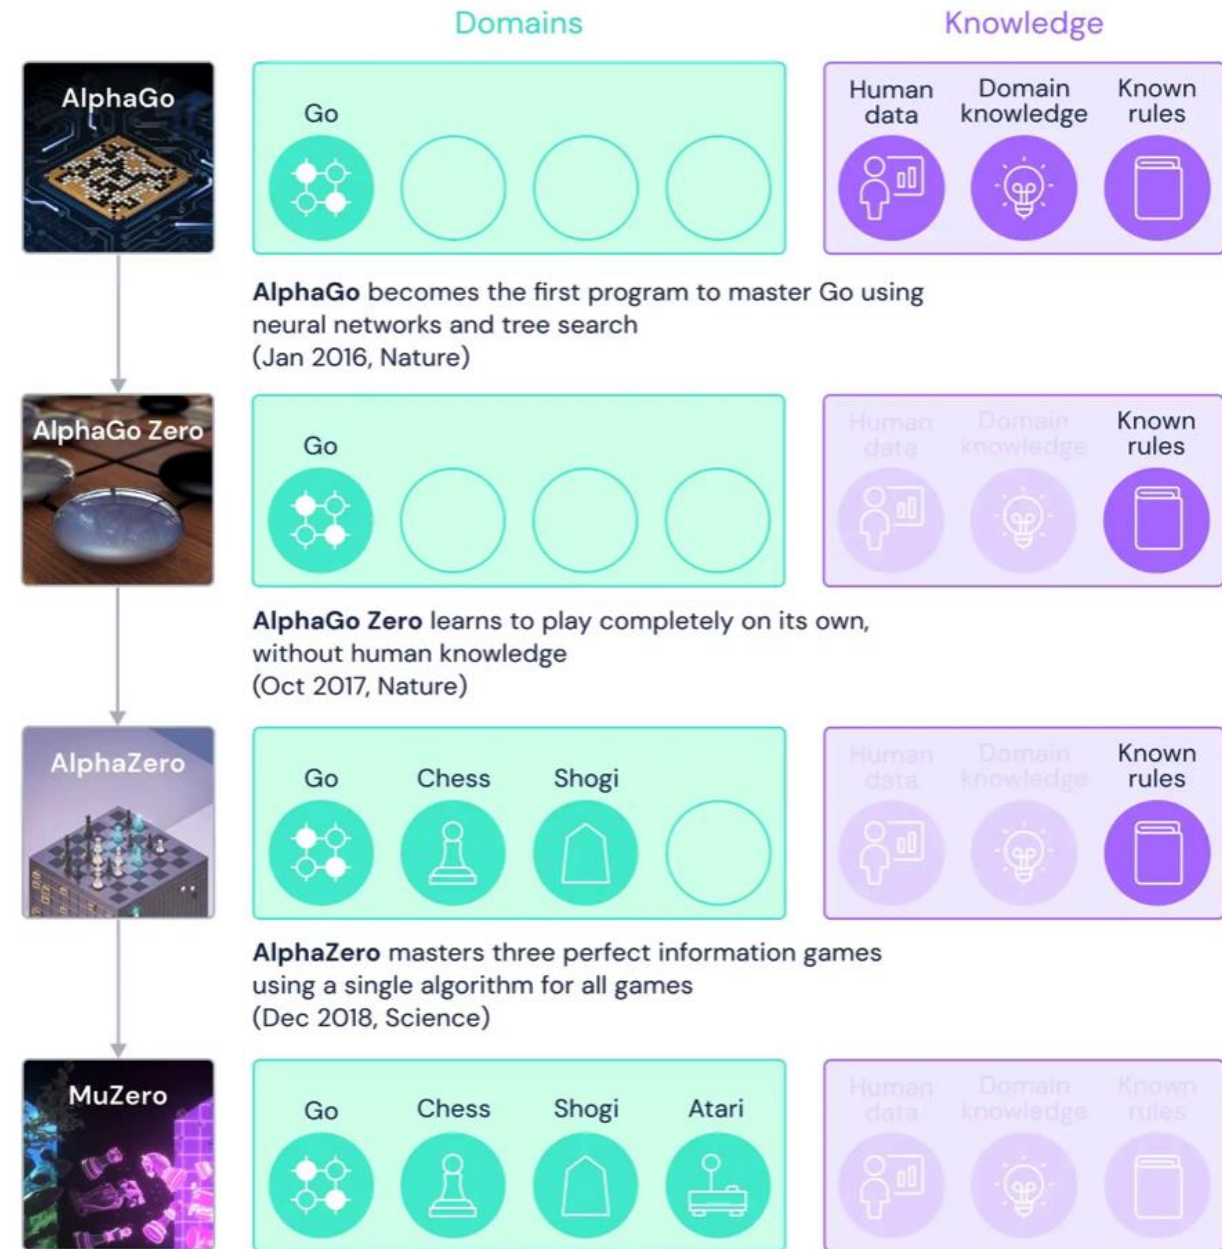
- AlphaGo [1]
- AlphaGo Zero [2]
- Alpha Zero [3]
- Mu Zero [4]

[1] Silver, David, et al. "Mastering the game of Go with deep neural networks and tree search." *nature*

[2] Silver, David, et al. "Mastering the game of go without human knowledge." *nature*

[3] Silver, David, et al. "A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play." *Science*

[4] Schrittwieser, Julian, et al. "Mastering atari, go, chess and shogi by planning with a learned model." *Nature*



<https://www.deepmind.com/blog/muzero-mastering-go-chess-shogi-and-atari-without-rules>

Useful applications

Deep reinforcement learning for the control of microbial co-cultures in bioreactors

Neythen J. Treloar, Alex J. H. Fedorec, Brian Ingalls , Chris P. Barnes 

Deep reinforcement learning for optimal experimental design in biology

Neythen J. Treloar , Nathan Braniff, Brian Ingalls, Chris P. Barnes 

Article | [Published: 09 June 2021](#)

A graph placement methodology for fast chip design

[Azalia Mirhoseini](#) , [Anna Goldie](#) , [Mustafa Yazgan](#), [Joe Wenjie Jiang](#), [Ebrahim Songhori](#), [Shen Wang](#), [Young-Joon Lee](#), [Eric Johnson](#), [Omkar Pathak](#), [Azade Nazi](#), [Jiwoo Pak](#), [Andy Tong](#), [Kavya Srinivasa](#), [William Hang](#), [Emre Tuncer](#), [Quoc V. Le](#), [James Laudon](#), [Richard Ho](#), [Roger Carpenter](#) & [Jeff Dean](#)

[Nature](#) **594**, 207–212 (2021) | [Cite this article](#)

Article | [Open Access](#) | [Published: 16 February 2022](#)


Magnetic control of tokamak plasmas through deep reinforcement learning

[Jonas Degraeve](#), [Federico Felici](#) , [Jonas Buchli](#) , [Michael Neunert](#), [Brendan Tracey](#) , [Francesco Carpanese](#), [Timo Ewalds](#), [Roland Hafner](#), [Abbas Abdolmaleki](#), [Diego de las Casas](#), [Craig Donner](#), [Leslie Fritz](#), [Cristian Galperti](#), [Andrea Huber](#), [James Keeling](#), [Maria Tsimpoukelli](#), [Jackie Kay](#), [Antoine Merle](#), [Jean-Marc Moret](#), [Seb Noury](#), [Federico Pesamosca](#), [David Pfau](#), [Olivier Sauter](#), [Cristian Sommariva](#), [Stefano Coda](#), [Basil Duval](#), [Ambrogio Fasoli](#), [Pushmeet Kohli](#), [Koray Kavukcuoglu](#), [Demis Hassabis](#) & [Martin Riedmiller](#) 


[Nature](#) **602**, 414–419 (2022) | [Cite this article](#)

Starting at the end

1. Multi armed bandits
2. Finite Markov decision processes
3. Dynamic programming
4. Reinforcement learning



Markov decision process



We work with discrete timesteps: $t_0, t_1, t_2, t_3, \dots$


In a fully observable MDP all information is encoded in the state i.e. the past is irrelevant:

- Most RL theory is based on this assumption

In a partially observable MPD (POMDP) only partial information can be observed and therefore history is important:

- We can include history using e.g. recurrent neural networks

Markov decision process




- Agent: yellow ball
- Environment: grid
- Action: move (U,D,L,R)
- State: current position
- Reward: reach the goal position and avoid penalty positions


| | | | | | | | | | |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0.00 ● | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↖ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | | | | | 0.00 ↕ | | | | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ |
| 0.00 ↖ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↕ | 0.00 ↖ |

Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: and introduction (2nd ed.). The MIT Press.
https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html

Markov decision process




- Agent: large green ball
- Environment: area
- Action: apply thrusters (U,D,L,R)
- State: position of agent, green and red ball
- Reward: +ve close to the small green ball, -ve close to the red ball




Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: and introduction (2nd ed.). The MIT Press.

<https://cs.stanford.edu/people/karpathy/reinforcejs/puckworld.html>

Markov decision process




- Agent: blue ball
- Environment: area
- Action: apply thrusters (U,D,L,R)
- State: Agent has 30 eye sensors which it uses to measure the range and velocity of sensed object. Agent also knows its own velocity
- Reward: +ve eat the red apples, -ve eat the green poison apples



Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: and introduction (2nd ed.). The MIT Press.

<https://cs.stanford.edu/people/karpathy/reinforcejs/waterworld.html>

Backup diagrams



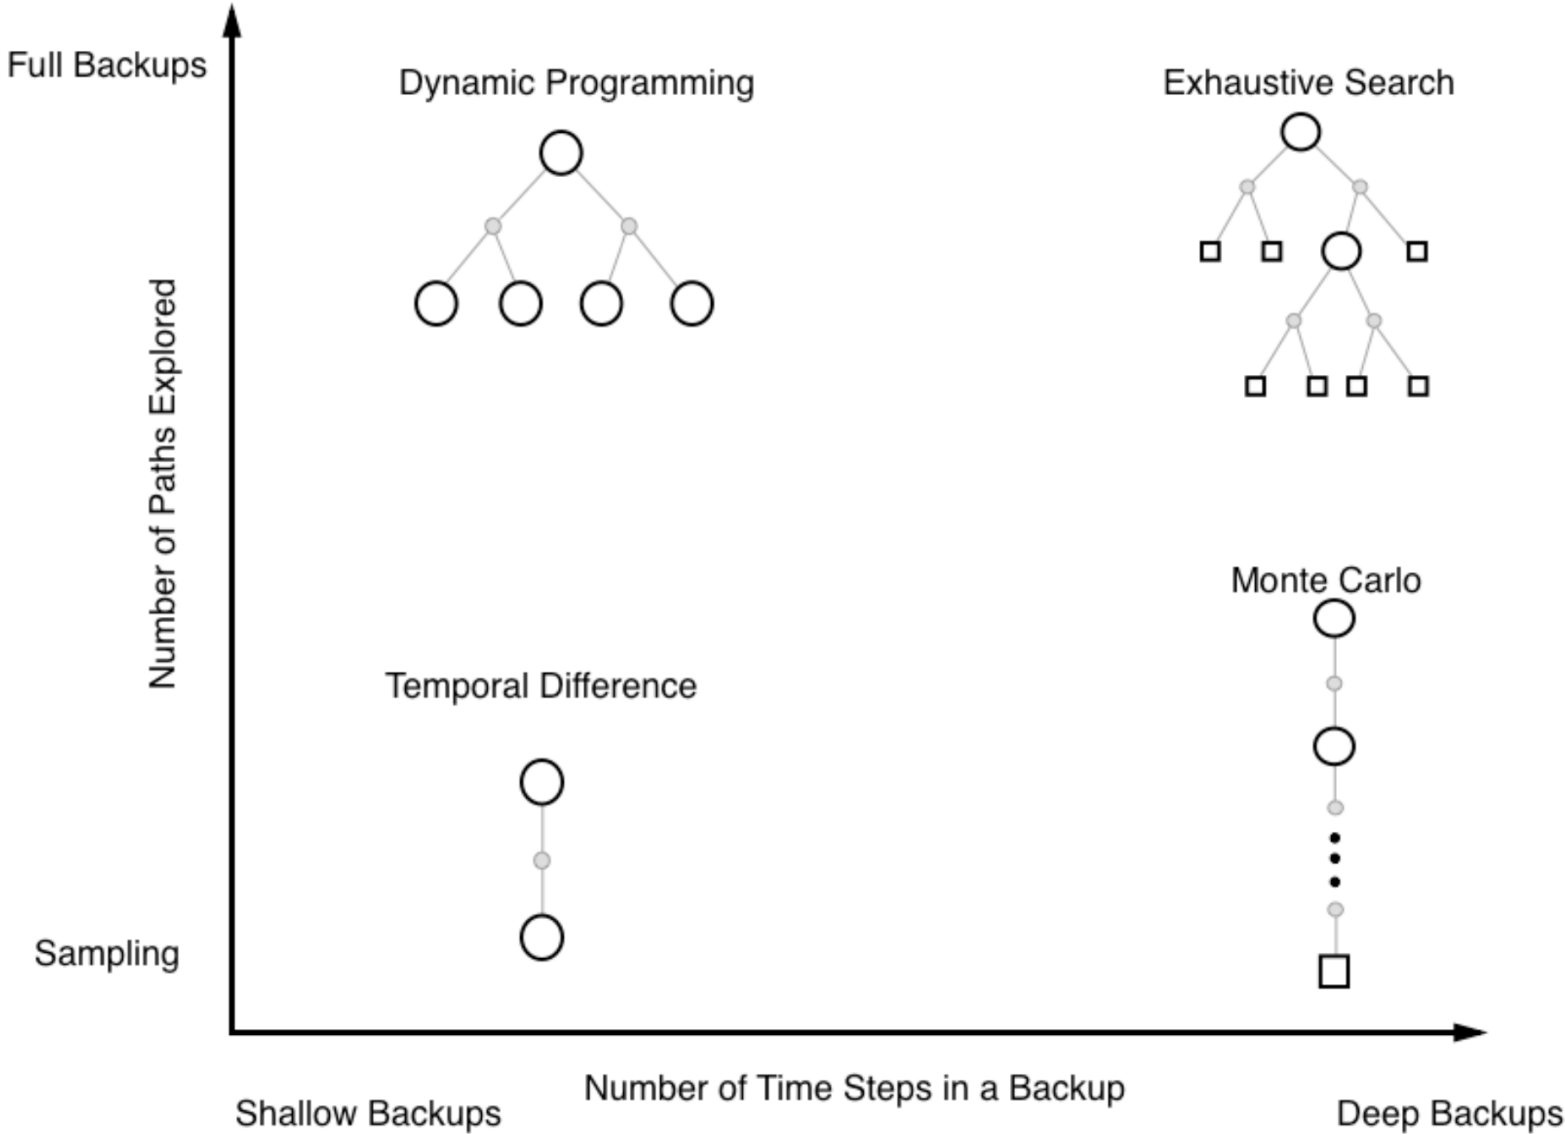
Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: and introduction (2nd ed.). The MIT Press.

Episodes

- We often train over repeated **episodes**
- An episode is a sequence of states, actions and rewards that continues until a terminal state is reached
- For example:
 - A level in a video game
 - An experiment
 - A match of a board game
- It is also possible to train on one infinitely long episode:
 - The puckworld and waterworld examples




Have to be model based

Can learn from experience
(these are RL methods)



Reinforcement learning

Agent learns an optimal behaviour **policy** by interacting with its **environment** to maximise its total **reward**, the **return**



The value function

- Many RL methods aim to learn a **value function**
- This is an estimate of the **expected discounted future return** the agent will receive and is learned through trial and error
- This can be learned for different states (a **state value function**) or different combinations of state and action (a **state-action value function**)

The state, value function $V(s)$

- In the grid world example, the state is the position on the grid
- For each state a **value** is learned, which is an estimate of the expected discounted return for an agent in that state.
- The table is a state value function $V(s)$
- What is the value of being in the top left corner? $s = (0,0)$
 - 0.22
- What are the highest and lowest value states?
 - (5,5) and (3,3)

| | | | | | | | | | |
|-----------|-----------|-----------|----------------------|-----------|----------------------|----------------------|-----------|----------------------|-----------|
| 0.22 ↗ | 0.25 ↗ | 0.27 ↗ | 0.31 ↗ | 0.34 ↗ | 0.38 ↓ | 0.34 ↖ | 0.31 ↖ | 0.34 ↗ | 0.38 ↓ |
| 0.25 → | 0.27 → | 0.31 → | 0.34 → | 0.38 → | 0.42 ↓ | 0.38 ← | 0.34 ↔ | 0.38 → | 0.42 ↓ |
| 0.27 ↑ | | | | | 0.46 ↓ | | | | 0.46 ↓ |
| 0.20 ↖ | 0.22 ↗ | 0.25 ↓ | -0.78 ↖ R -1.0 | | 0.52 → | 0.57 → | 0.64 ↓ | 0.57 ↖ | 0.52 ↖ |
| 0.22 ↗ | 0.25 ↗ | 0.27 ↓ | 0.25 ↖ | | 0.08 ↓ R -1.0 | -0.36 → R -1.0 | 0.71 ↓ | 0.64 ← | 0.57 ← |
| 0.25 ↖ | 0.27 ↖ | 0.31 ↓ | 0.27 ↖ | | 1.20 ↖ R 1.0 | 0.08 ← R -1.0 | 0.79 ↓ | -0.29 ← R -1.0 | 0.52 ↓ |
| 0.27 ↖ | 0.31 ↖ | 0.34 ↓ | 0.31 ← | | 1.08 ↑ | 0.97 ← | 0.87 ← | -0.21 ← R -1.0 | 0.57 ↓ |
| 0.31 ↖ | 0.34 ↖ | 0.38 ↓ | -0.58 ↓ R -1.0 | | -0.03 ↑ R -1.0 | -0.13 ↑ R -1.0 | 0.79 ↑ | 0.71 ← | 0.64 ← |
| 0.34 → | 0.38 → | 0.42 → | 0.46 → | 0.52 → | 0.57 → | 0.64 → | 0.71 ↑ | 0.64 ↖ | 0.57 ↖ |
| 0.31 ↖ | 0.34 ↖ | 0.38 ↖ | 0.42 ↖ | 0.46 ↖ | 0.52 ↖ | 0.57 ↖ | 0.64 ↖ | 0.57 ↖ | 0.52 ↖ |

The state-action, value function $Q(s, a)$

- The state-action value function $Q(s, a)$, this is the value obtained by starting in state s and taking action a
- In grid world we can choose from up to four actions, U,D,L,R
- E.g. for the top left corner $s = (0,0)$, we have two available actions, R or D. $Q((0,0), R) = Q((0,0), D) = 0.25$. Because starting from $(0,0)$ and moving either R or D we end up in states with value 0.25
- What is $Q((3,2), R)$?
 - -0.78
- What is $Q((6,5), U)$?
 - 1.2
- Learning this is the basis of many basic and advanced RL techniques

| | | | | | | | | | |
|-----------|-----------|-----------|---------------------|-----------|---------------------|---------------------|-----------|---------------------|-----------|
| 0.22 ↗ | 0.25 ↗ | 0.27 ↗ | 0.31 ↗ | 0.34 ↗ | 0.38 ↓ | 0.34 ↖ | 0.31 ↖ | 0.34 ↗ | 0.38 ↓ |
| 0.25 → | 0.27 → | 0.31 → | 0.34 → | 0.38 → | 0.42 ↓ | 0.38 ← | 0.34 ↔ | 0.38 → | 0.42 ↓ |
| 0.20 ↑ | | | | | 0.46 ↓ | | | | 0.46 ↓ |
| 0.20 ↙ | 0.22 ↗ | 0.25 ↓ | -0.78 ↖ R-1.0 | | 0.52 → | 0.57 → | 0.64 ↓ | 0.57 ↖ | 0.52 ↖ |
| 0.22 ↗ | 0.25 ↗ | 0.27 ↓ | 0.25 ↖ | | 0.08 ↓ R-1.0 | -0.36 → R-1.0 | 0.71 ↓ | 0.64 ← | 0.57 ← |
| 0.25 ↗ | 0.27 ↗ | 0.31 ↓ | 0.27 ↖ | | 1.20 ↖ R-1.0 | 0.08 ← R-1.0 | 0.79 ↓ | -0.29 ← R-1.0 | 0.52 ↓ |
| 0.27 ↗ | 0.31 ↗ | 0.34 ↓ | 0.31 ← | | 1.00 ↑ | 0.97 ← | 0.87 ← | -0.21 ← R-1.0 | 0.57 ↓ |
| 0.31 ↗ | 0.34 ↗ | 0.38 ↓ | -0.58 ↓ R-1.0 | | -0.03 ↑ R-1.0 | -0.13 ↑ R-1.0 | 0.70 ↑ | 0.71 ← | 0.64 ← |
| 0.34 → | 0.38 → | 0.42 → | 0.46 → | 0.52 → | 0.57 → | 0.64 → | 0.71 ↑ | 0.64 ↖ | 0.57 ↖ |
| 0.31 ↙ | 0.34 ↙ | 0.38 ↙ | 0.42 ↙ | 0.46 ↙ | 0.52 ↙ | 0.57 ↙ | 0.64 ↙ | 0.57 ↙ | 0.52 ↙ |

Discounting

Agent learns from experience to maximise its total reward, the **return**. During learning this is discounted by a factor γ between 0 and 1


- $\gamma = 1$ a reward in the future is worth just as much as a reward now. This removes the time agency of the agent, a longer path to the reward will be just as optimal as a shorter path to the same reward. Can lead to instability, especially when episodes not time constrained
- $\gamma = 0$ agent is only concerned about the next reward, will forgo large future rewards just to get a small reward now.
- 0.95 is a good starting point
- What is the discount factor used to learn this value function?
 - 0.9, because the value decays by a factor of 0.9 for each step away from the goal

$$G = \sum_{t=0} \gamma^t r_t = r_0 + \gamma r_1 + \gamma^2 r_2 + \gamma^3 r_3 + \gamma^4 r_4 + \dots$$

$$\gamma = 0: G = r_0$$

$$\gamma = 1: G = r_0 + r_1 + r_2 + r_3 + r_4 + \dots$$

$$\gamma = 0.95: G = r_0 + 0.95 * r_1 + 0.9 * r_2 + 0.86 * r_3 + 0.81 * r_4 + \dots$$



This is why the value of each position decreases as we move away from the goal

The policy


- The agent's overall goal is to learn the optimal behaviour **policy**
- The policy is the learned decision-making process the agent uses to select the next action, a , based on the current state, s
- The greedy policy is found by choosing the action with highest value for each state
- In grid world the learned policy is represented by the arrows
- What route does the agent take from the initial state (0,0) to the goal state (5,5) under this policy?



Terminology review

Agent learns an optimal behaviour **policy** by interacting with its **environment** to maximise its total **reward**, the **return**

A lot of RL methods learn a state-action **value function** $Q(s, a)$. This is a learned estimate of the **expected discounted return** received after visiting every state-action pair




Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: and introduction (2nd ed.). The MIT Press.

Exploration

- A key concept in RL is the exploration/exploitation trade off
- We need to explore so that we can discover new behaviours which might be optimal
- But we don't want to explore too much otherwise we will never exploit the knowledge we have gained
- A commonly used exploration policy is **ϵ -greedy**:
 - The **explore rate**, ϵ , is the probability that the agent takes a random action
 - With probability $1 - \epsilon$ the agent will take the greedy action $a = \max_a Q(s, a)$
 - The explore rate can be set to a constant small value (often 0.05) or can be set to 1 initially and decay throughout training

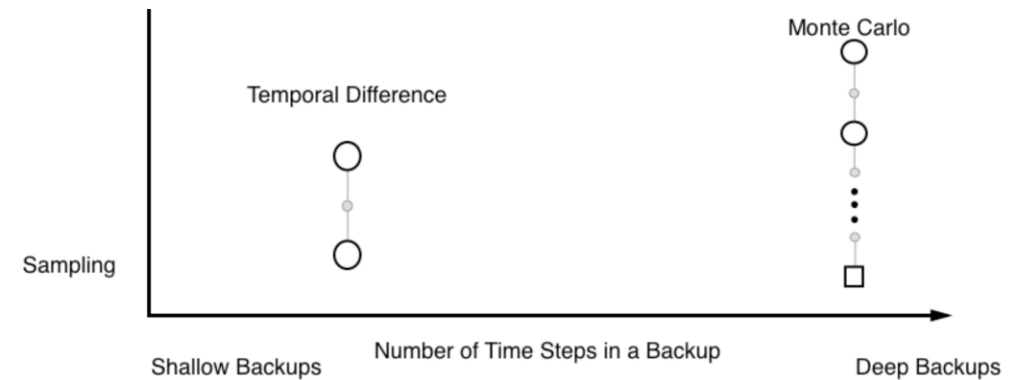
On vs Off policy methods

- Often, we use the ϵ -greedy policy to explore, but we want to learn the optimal greedy policy
- Meaning our **behaviour policy** (ϵ -greedy) is different to our **target policy** (greedy)
- If our behaviour policy is the same as our target policy this is called an **on-policy** method e.g. SARSA
- If our behaviour policy is different from our target policy this is called an **off-policy** method e.g. Q-learning
- When could it be advantageous to choose each type of method?



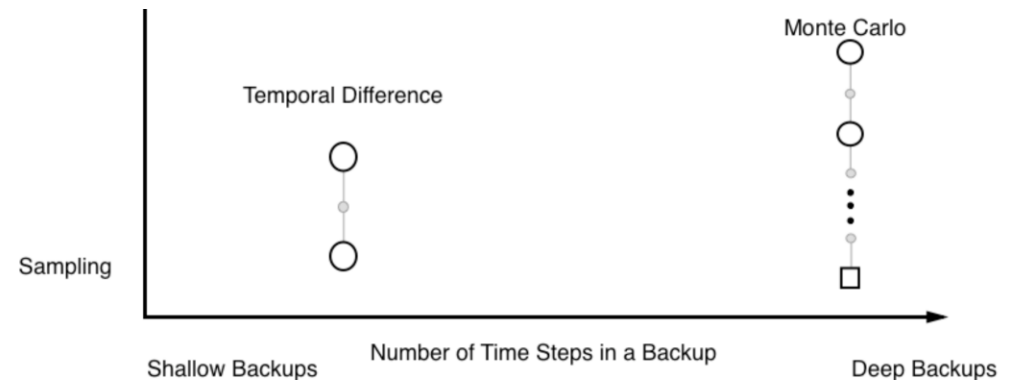
Learning the value function: Monte Carlo

- Now all we need to do is learn the $Q(s, a)$ function
- In Monte Carlo RL we learn the mean return for each s, a pair
 - Keep track of sum of the return $S_G(s, a)$
 - And a count of how many times each s, a pair visited $C(s, a)$
 - So that $Q(s, a) = S_G(s, a) / C(s, a)$
- Implementation is a bit tricky, but principle is simple
- $G(s_t, a_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \gamma^T r_{t+T}$
- $Q(s_t, a_t) = S_G(s_t, a_t) / C(s_t, a_t)$
- We are restricted to updating only after an episode has finished, because we need to know the full sequence of states, actions and rewards
 - Won't work for non-episodic tasks and can cause problems even for episodic ones
- Off policy learning is more difficult (importance sampling in Sutton and Barto)



Temporal difference methods

- Temporal difference methods replace the observed future return with an estimate of the future return from the current value function $Q(s_t, a_t)$
 - This is called bootstrapping
 - This means we don't have to wait until the end of the episode to learn
 - Much easier to do off-policy learning



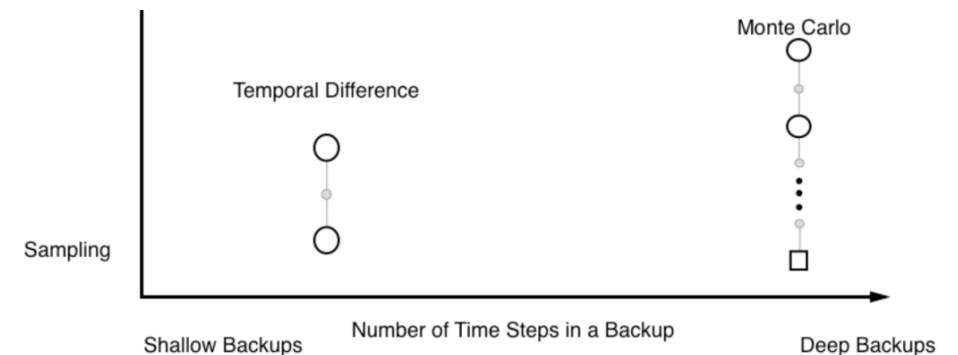
Learning the value function: on policy temporal difference (SARSA)

- Instead of using observed return:
 - $r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \gamma^T r_{t+T}$
- We can bootstrap off current $Q(s, a)$ function to estimate the expected future return:

- $r_t + \gamma Q(s_{t+1}, a_{t+1})$

Reward from time = t Expected return from time $t + 1$ onwards

- We iteratively update the value function with learning rate α
 - $Q(s_t, a_t) = Q(s_t, a_t) + \alpha (r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$

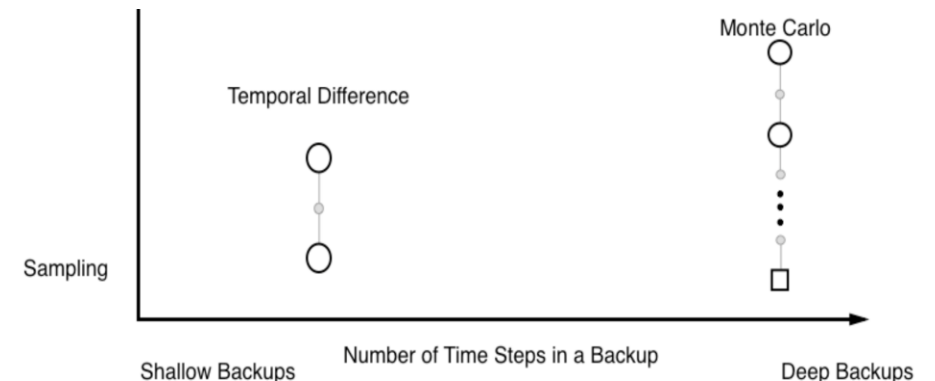


Learning the value function: off policy temporal difference (Q-learning)

- Instead of using observed return:
 - $r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \gamma^T r_{t+T}$
- We can bootstrap off current $Q(s, a)$ function to estimate the expected future return:
 - $r_t + \gamma \max_a Q(s_{t+1}, a)$

Reward from time = t Expected return from time $t + 1$ onwards

- We iteratively update the value function with learning rate α
 - $Q(s_t, a_t) = Q(s_t, a_t) + \alpha (r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$



Q-learning vs SARSA

SARSA (on policy)


- $$Q(s_t, a_t) = Q(s_t, a_t) + \alpha (r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

Update according to the action
taken by the behaviour policy

Q-learning (off policy)

- $$Q(s_t, a_t) = Q(s_t, a_t) + \alpha (r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$$

Update according to the greedy policy



Transitions

- $Q(s_t, a_t) = Q(s_t, a_t) + \alpha \left(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right)$
 - Q learning transition: (s_t, a_t, r_t, s_{t+1})
- $Q(s_t, a_t) = Q(s_t, a_t) + \alpha (r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
 - SARSA transition: $(s_t, a_t, r_t, s_{t+1}, a_{t+1})$
- $Q(s_t, a_t) = S_G(s_t, a_t) / C(s_t, a_t)$
 - Monte Carlo learns episodically so uses a sequence of transitions:
 - $(s_1, a_1, r_1), (s_2, a_2, r_2), \dots, (s_N, a_N, r_N)$

Pseudocode overview – a TD learning agent

Q learning agent:

- Q_function
 - This is the function that maps state-action pairs to learned values
 - In theory can be any function, we will look at using lookup tables and neural networks
- Policy
 - The agents policy maps a state observed from the environment to an action to take
 - We will use epsilon greedy policies to incorporate exploration
- Update_Q_function
 - This function takes a transition and updates our current value function using the Q learning or SARSA update rules

Pseudocode overview – training loop

For each episode:

 while the episode has not reached a terminal state:

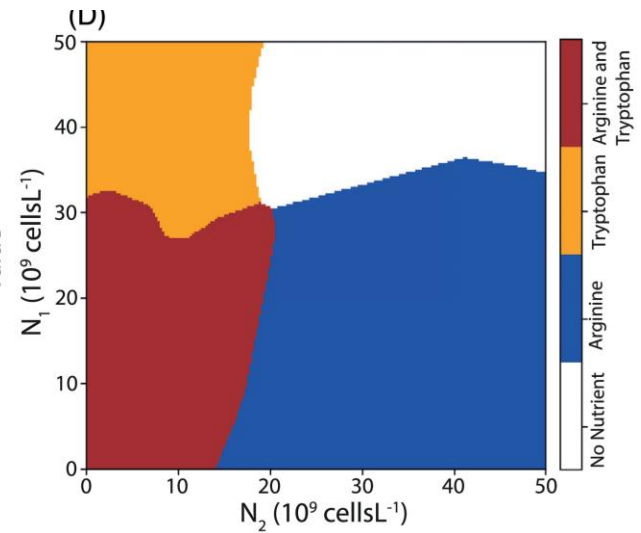


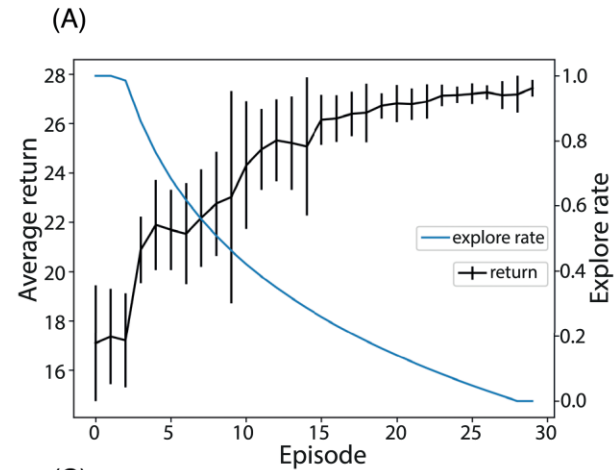

 get an action from the agent's policy

 apply that action to the environment to update the current state

 if we are training a temporal difference agent update the value function

 if we are training a Monte Carlo agent update the value function

Example training output



Deep reinforcement learning for the control of microbial co-cultures in bioreactors

Neythen J. Treloar, Alex J. H. Fedorec, Brian Ingalls, Chris P. Barnes

Summary

- RL can be used to solve control problems
- Often this is done by learning a value function from a sequence of states, actions and rewards
- This can be done episodically using Monte Carlo techniques
- Or online using temporal difference
 - Q-learning (off policy)
 - SARSA (on policy)
- Temporal difference methods allow us to update the value function during an episode and easily do off policy learning